

# On Julia’s efficient algorithm for subtyping union types and covariant tuples (Pearl)

Anonymous Authors

Anonymous Affiliation

## Abstract

We describe the algorithm implemented in the Julia programming language runtime to decide subtyping on a simple type system with union types, covariant tuples, and literals. This algorithm is immune from the space-explosion and expressiveness problems of standard algorithms that normalise types into disjunctive normal form ahead-of-time. We prove this algorithm correct and complete against a semantic-subtyping denotational model in Coq.

**2012 ACM Subject Classification** Theory of computation → Type theory

**Keywords and phrases** Type systems, Subtyping, Union types

## 1 Subtyping union types and covariant tuples

Union types are increasingly common in mainstream languages. In some cases, as Julia [2] or TypeScript [4], they are exposed at the source level. In others, such as Hack [1], they are only used internally when performing type inference. In all cases, they play a key role in the expressiveness of the language.

Rules for subtyping union types have been known for a long time. Following Vouillon [6] they may be written as:

$$\frac{t_1 <: t \quad t_2 <: t}{\text{Union}\{t_1, t_2\} <: t} \qquad \frac{t <: t_1}{t <: \text{Union}\{t_1, t_2\}} \qquad \frac{t <: t_2}{t <: \text{Union}\{t_1, t_2\}}$$

These rules are asymmetrical. Following from set-theoretic intuitions, and as made formal by semantic subtyping models, when a union type appears on the left-hand side of a subtype judgment, then *all* its components must be subtypes of the right-hand side; when a union type appears on the right-hand side of a subtype judgment, then there must *exist* a component that is a supertype of the left-hand side.

It has also been known for a long time that, in the presence of covariant tuples, the above rules are not complete with respect to a semantic subtyping model [5]. Let us recall the subtype rule for covariant tuples:

$$\frac{t_1 <: t'_1 \quad t_2 <: t'_2}{\text{Tuple}\{t_1, t_2\} <: \text{Tuple}\{t'_1, t'_2\}}$$

In a semantic subtyping model, covariant types should be *distributive* with respect to unions; that is, the following (and the opposite) derivation should hold:

$$\text{Tuple}\{\text{Union}\{t_1, t_2\}, t\} <: \text{Union}\{\text{Tuple}\{t_1, t\}, \text{Tuple}\{t_2, t\}\}$$

Here the rule for tuples cannot be applied, and a derivation must immediately pick either the first or second component of the union type; as a result, it is impossible to complete the derivation. The standard approach to solve this conundrum is to *rewrite all types into their disjunctive normal form* (DNF), that is as unions of union-free types, *before* building the derivation. The correctness of this rewriting step is justified by the semantic-subtyping denotational model, as in [3], and the resulting subtype algorithm can be proved both correct

and complete. However, this standard algorithm based on ahead-of-time normalization has *two major drawbacks*.

The first drawback is that the normalization phase rewrites types into potentially *exponentially bigger* types. This is a problem in practice. Previous work [7] instrumented Julia's subtype decision procedure and when executing library code, routinely observed queries involving types as the one below:

```

42 Tuple{Tuple{Union{Int64, Bool}, Union{String, Bool}, Union{String, Bool},
43           Union{String, Bool}, Union{Int64, Bool}, Union{String, Bool},
44           Union{String, Bool}, Union{String, Bool}, Union{String, Bool},
45           Union{String, Bool}, Union{String, Bool}, Union{String, Bool},
46           Union{String, Bool}, Union{String, Bool}, Union{String, Bool}}, Int64}

```

Normalizing this type before attempting to decide subtyping is not a realistic option. In practice, such types are often matched against structurally simpler types like

```

49 Tuple{Tuple{Any, Any, Any, Any, Any, Any, Any, Any, Any, Any, Any, Any, Any, Any, Any},
50         Any}

```

(where in Julia `Any` is the supertype of all types); in these cases, it is possible to decide subtyping in linear time without ahead-of-time normalisation.

The second drawback of the ahead-of-time normalisation phase is that it does not interact well with other type system features. For instance, *invariant constructors* make lifting all union types to the top level unsound. Consider the type `Array{Union{ $t_1$ ,  $t_2$ }}`, where `Array` is an invariant unary constructor. This type denotes the set of arrays whose elements are either of type  $t_1$  or  $t_2$ ; it would be incorrect to rewrite it as `Union{Array{ $t_1$ }, Array{ $t_2$ }}`, as this latter type denotes the set of arrays whose elements are either all of type  $t_1$  or all of type  $t_2$ . A weaker disjunctive normal form, only lifting union types inside each invariant constructor, can circumvent this problem. However, doing so only reveals a deeper problem in the presence of both invariant constructors and *existential types*. Consider the judgment below:

```

63 Array{Union{Tuple{ $t_1$ }, Tuple{ $t_2$ }}} <:  $\exists T$ . Array{Tuple{T}}

```

This judgment holds by taking  $T = \text{Union}\{t_1, t_2\}$ . Since all types are in weak normal form, a standard algorithm would initially perform the left-to-right subtype check due to the outer invariant constructor. This step would generate the constraint  $T >: \text{Union}\{t_1, t_2\}$ . As a consequence, the subsequent right-to-left check due to the invariant constructor fails. Indeed this requires proving that `Tuple{T} <: Union{Tuple{ $t_1$ }, Tuple{ $t_2$ }}`, which in turns attempts to prove either  $T <: t_1$  or  $T <: t_2$ , both unprovable under the assumption that  $T >: \text{Union}\{t_1, t_2\}$ . The key to derive a successful judgment for this relation is to rewrite the right-to-left check into `Tuple{T} <: Tuple{Union{ $t_1$ ,  $t_2$ }}`, which is provable. This is a sort of *anti-normalisation* rewriting that must be performed on sub-judgments of the derivation, and to the best of our knowledge it is not part of any subtype algorithm based on ahead-of-time disjunctive normalisation.

In this *pearl paper* we describe the keys ideas used by the subtype algorithm implemented in the Julia language to deal with union types and covariant tuples, and we will argue that these avoid the two drawbacks above. To avoid being drawn in the vast complexity of Julia type algebra, we focus on a minimal language featuring union types, covariant tuples, and literals. This tiny language is expressive enough to highlight the decision strategy, and make this implementation technique known to a wider audience. While Julia implementation shows that this technique extends, among others, to invariant constructors and existential types [7],

we expect that it can be leveraged in many other modern language designs. Additionally we prove in Coq that the algorithm is correct and complete with respect to a standard semantics subtyping model.

## 2 A space-efficient algorithm

The greatest difficulty of subtyping union types is searching the entire space of possible forall/exist quantifications of union types. We have seen that syntax-directed rules need ahead-of-time normalisation in the presence of covariant tuples.

The algorithm we propose for deciding subtyping works by iterating through choices made at unions along paths complete through the two types. In this section, we will describe the implementation of the algorithm in OCaml and will present our proof of correctness in the following section.

We focus on a core type language composed of binary unions, covariant binary tuples, and atom types:

```

type typ =
  | Atom of int
  | Tuple of typ * typ
  | Union of typ * typ

```

Atom types are singletons with respect to subtyping, and are ranged over by  $A, B, \dots$ . Abusing the notation, we will assume that we also have unary tuples; these can be easily encoded.

Stack		Type		Base Query
$\forall$	$\exists$	$\forall$	$\exists$	
L	L			$A <: A$
L	R			$A \not<: B$
R	L			$B \not<: A$
R	R			$B <: B$

$\text{Tuple}\{\text{Union}\{A, B\}\} <: \text{Union}\{\text{Tuple}\{A\}, \text{Tuple}\{B\}\}$

Figure 1 Subtyping decision procedure example

As a running example we use the distributivity example from the introduction:

$$\text{Tuple}\{\text{Union}\{A, B\}\} <: \text{Union}\{\text{Tuple}\{A\}, \text{Tuple}\{B\}\}$$

In Figure 1 we sketch the corresponding execution of the subtyping decision procedure. Recall that to check if the relation holds, the algorithm needs to ensure that for every choice that could be made at a union types on the left hand side of the judgment there is a set of choices for the union types on the right hand side such that subtyping holds. In this example, the relation holds: no matter if we choose  $A$  or  $B$  on the left hand side inside the covariant tuple, we can always pick a matching type on the right hand side. The key challenge for a subtyping algorithm not based on ahead-of-time normalisation is to enumerate every possible choice of the components of the unions appearing in the left and right hand sides of the judgments.

To do this, we base our approach on iterators. The `st_choice` type represents whether the algorithm takes the left (first) or the right (second) component of a union type:

```
type st_choice = Left | Right
```

The algorithm stores the choice made at each union at the present iteration in *two stacks of st\_choices*, one for each side of the subtype relation. Each `st_choice` stack can be seen as the state of an iterator that enumerates every union-type induced alternative in a given type. One stack will be used to ensure that all *forall* choices have been explored; the other stack to search for a matching *exist* choice. In figure 1, these are depicted as the  $\forall$  and  $\exists$  stacks.

It is possible to define both an initial element and next state function for iterators defined using `st_choice` stacks.

```
let rec initial_choice (a:typ) =
  match a with
  | Atom i -> []
  | Tuple(t1,t2) -> (initial_choice t1) @ (initial_choice t2)
  | Union(l,r) -> Left::(initial_choice l)
```

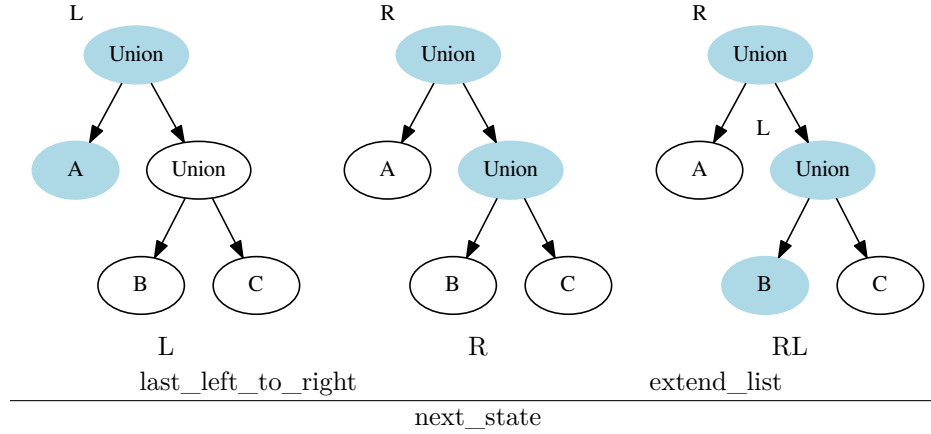
We choose the initial value for our iterators to be the state where, every time a union type is encountered in the structure of the type, the iterator takes the `Left` choice (as implemented in `initial_choice` – as we shall see when we will discuss the `extend_list` function, empty lists are expanded on demand with the appropriate `Left` choices).

Given this definition for the initial state, we can then define the iteration function. We break the iteration function, `next_state`, up into two helper functions. The operation of these helpers is illustrated in figure 2. Here, starting from the initial state (where the iterator takes the `Left` choice at the root of the type), `last_left_to_right` converts the final `Left` choice into a `Right` choice and truncates the remaining choices.

```
let rec last_left_to_right
  (acc:st_choice list) (ll : st_choice list option) = function
  | Left::tl -> last_left_to_right (Left::acc) (Some (Right::acc)) tl
  | Right::tl -> last_left_to_right (Right::acc) ll tl
  | [] -> option_map List.rev ll
```

For example, given the type `Union{Union{A, B}, C}` and the choice list `LL`, `last_left_to_right` will produce the choice list `LR` (`Left` and `Right` are shortened here with `L` and `R`); if given `LR`, it will produce `R` (as the final left choice is at the top-level and it truncates the remainder of the choice list).

In Figure 2, `last_left_to_right` produced a partial path, as the right child of the root union is a union itself. To solve this, `extend_list` finds where the new choice list runs out and fills it out with left choices to be valid with respect to the type, producing a valid state. `extend_list` pads the list out to take the left choice at this child union, returning the path to validity.



■ **Figure 2** State-stepping operation for choice lists

```

147 let rec extend_list (a:typ) (ls:st_choice list) = match (a,ls) with
148   | (Atom i, _) -> ([], ls)
149   | (Tuple(t1,t2), _) -> let (hd,t1) = extend_list t1 ls in
150                           let (hd2,t12) = extend_list t2 t1 in
151                           (hd @ hd2, t12)
152   | (Union(l,r), Left::rs) -> let (hd,t1) = extend_list l rs in (Left::hd,t1)
153   | (Union(l,r), Right::rs) -> let (hd,t1) = extend_list r rs in (Right::hd,t1)
154   | (Union(l,r), []) -> (Left::initial_choice l,[])

```

Finally, the `next_state` function combines these operations into a single step operation. It will take the deepest alternative choice at a union using `last_left_to_right`, then add however many left choices are required to make the path valid using `extend_list`. The complete `next_state` operation shown in figure 2 takes the choice stack `L` and the type `Union{A, Union{B, C}}` and produces the choice stack `RL`. To do so, it takes the left choice, converts it to a right, then pads the list out with lefts until the path is valid with respect to the type.

```

162 let rec next_state (a:typ) (ls:st_choice list) =
163   option_map fst (option_map (extend_list a) (last_left_to_right [] None ls))

```

Now that we have defined the core iterator infrastructure, we show how it is used to decide subtyping queries. The algorithm proceeds by maintaining two iterators (one for each side of the subtyping judgment) over choices-at-unions, checking that for every instantiation of the left hand type there exists an instantiation of the right hand type. We will first define the fundamental subtyping relation, used to decide subtyping relationships when given instantiations of the left and right hand types, then describe the algorithm that iterates through those types.

```

171 type st_res =
172   | Subtype of st_choice list * st_choice list
173   | NotSubtype

```

The type `st_res` represents the results of a single base subtyping query. The query can either succeed, in which case it provides the unused portion of the input stacks, or fail.

```

176 let rec base_subtype (a:typ) (b:typ) (fa:st_choice list) (ex : st_choice list)

```

```

177     match (a,b,fa,ex) with
178     | (Atom i, Atom j, _, _) -> if i == j then Subtype(fa, ex) else NotSubtype
179     | (Tuple(ta1, ta2), Tuple(tb1, tb2), _, _) ->
180         (match base_subtype ta1 tb1 fa ex with
181         | Subtype(cfa, cex) -> base_subtype ta2 tb2 cfa cex
182         | NotSubtype -> NotSubtype)
183     | (Union(a1,a2),b,choice::fa,ex) ->
184         base_subtype (match choice with Left -> a1 | Right -> a2) b fa ex
185     | (a,Union(b1,b2),fa,choice::ex) ->
186         base_subtype a (match choice with Left -> b1 | Right -> b2) fa ex

```

The function `base_subtype` is responsible for using two paths—one for the left and right hand sides of the subtyping relation—and checking that the subtype relation holds with respect to those paths. Given a basic type—an atom or a tuple—it will either check equality or recur respectively. Given a union type, it will choose the component of the union type following the choice stacks and continue recursively. The function then returns an instance of `st_res`, which gives the remaining choice stacks if successful or nothing otherwise.

Finally, to check subtyping, we need to iterate through both types in the correct quantification order. Subtyping holds if, for every instantiation of the left there exists an instantiation of the right hand side of the subtyping judgment such that the subtype relation holds. To check this, we use a simple brute force approach sitting atop the choice stack iterator infrastructure described previously. The algorithm is implemented in two functions. `ex_subtype` checks that there exists an instantiation of the right-hand-side for a given left-hand-side. `fa_ex_subtype` uses `ex_subtype` to check that for every instantiation of the LHS, there is an instantiation of the RHS such that the subtype relation holds.

```

201     let rec ex_subtype (a:typ)(b:typ)(fa:st_choice list)(cex : st_choice list) =
202         match base_subtype a b fa cex with (* is a <: b *)
203         | Subtype(_,_) -> true (* there exists a subtype *)
204         | NotSubtype ->
205             (match next_state b cex with (* step exists-env *)
206             | Some ns -> ex_subtype a b fa ns (* continue *)
207             | None -> false) (* no subtype; exit *)

```

If the current choice — given by `cex` — is a supertype of the given `a` according to `base_subtype`, then `ex_subtype` has found a valid instantiation of `b`. Therefore, there exists an instantiation of `b` that is a supertype of `a` and the result should be true. Otherwise, `ex_subtype` will use the iteration operation, `next_state`, to continue iterating through choices for `b` until it either finds an instantiation that is a supertype or runs out of instantiations.

```

214     let rec fa_ex_subtype (a:typ)(b:typ)(cfa:st_choice list) =
215         match ex_subtype a b cfa (initial_choice b) with (* a <: b wrt path? *)
216         | true -> (match next_state a cfa with
217         | Some ns -> fa_ex_subtype a b ns (* continue *)
218         | None -> true) (* all subtypes; is subtype *)
219         | false -> false (* exists a non-subtype; not subtype *)

```

The function `fa_ex_subtype` is similar; it checks that for every instantiation of `a`, there exists an instantiation of `b` such that subtyping holds. Checking for an instantiation of `b` is done using `ex_subtype`, while `fa_ex_subtype` maintains an iterator for `a`.

The operation of `fa_ex_subtype` and `ex_subtype` as well as their calls to `base_subtype` can be seen in figure 1. In the  $\forall$  column, we see the current state maintained by `fa_ex_subtype`.

225 The  $\exists$  column shows the state maintained by `ex_subtype` as it is called by `fa_ex_subtype`.  
 226 In the example, for a forall-list of L, we can find an exists-list of L such that `base_subtype`  
 227 holds. Similarly, for a forall-list of R, we can find an exists-list instantiation of R such that  
 228 subtyping holds. Therefore, for every instantiation of the left-hand-side, there exists an  
 229 instantiation of the right-hand-side such that subtyping holds and subtyping holds for the  
 230 type as a whole.

```
231 let rec subtype (a:typ) (b:typ) = fa_ex_subtype a b (initial_choice a)
```

232 Finally, `subtype` checks if `a` is a subtype of `b`. It seeds `fa_ex_subtype` with the initial  
 233 choice for `a`'s iterator, which then checks if for every instantiation of `a` there exists an  
 234 instantiation of `b` such that subtyping holds.

235 We have presented our subtyping algorithm using lists of choices. In a practical implemen-  
 236 tation, however, these lists of choices can be efficiently implemented (without allocation) by  
 237 means of bit sets. This is the approach taken in the Julia implementation of this algorithm.  
 238 With this optimization, the needed memory to decide a subtyping judgment is linear in the  
 239 total number of unions in the given types; the algorithm needs no allocation beyond that of  
 240 the choice stacks themselves.

### 241 3 Correctness and completeness

242 To prove correctness of our algorithm, we begin by formally specifying correctness for  
 243 subtyping. We then show that two subtyping algorithms— based on structural iterators and  
 244 choice lists—are correct with respect to this definition.

245 We base our definition of subtyping on a denotational semantics for types. We reduce  
 246 types in the type language including unions to sets of types in the type language without  
 247 unions through a simple transformation.

$$\begin{aligned} 248 \quad \llbracket A \rrbracket &= \{A\} \\ 249 \quad \llbracket \text{Union}\{t_1, t_2\} \rrbracket &= \llbracket t_1 \rrbracket \cup \llbracket t_2 \rrbracket \\ 250 \quad \llbracket \text{Tuple}\{t_1, t_2\} \rrbracket &= \{\text{Tuple}\{t'_1, t'_2\} \mid t'_1 \in \llbracket t_1 \rrbracket, t'_2 \in \llbracket t_2 \rrbracket\} \end{aligned}$$

251

253 Using this denotational semantics for types-with-unions, we can define subtyping as  
 254 if  $\llbracket t_1 \rrbracket \subseteq \llbracket t_2 \rrbracket$ , then  $t_1 <: t_2$ . Equivalently, we can state this as definition 1, which is  
 255 canonicalized in our Coq proof as the `NormalSubtype` relation.

256 ► **Definition 1** (Subtyping Correctness). *A subtyping relation  $<:$  is correct if  $t_1 <: t_2$  iff*  
 257  *$\forall t'_1 \in \llbracket t_1 \rrbracket, \exists t'_2 \in \llbracket t_2 \rrbracket, t_1 <: t_2$ .*

258 Proving a subtyping algorithm sound and complete is therefore equivalent to producing a  
 259 function of type `forall t1 t2:typ, {NormalSubtype t1 t2} + {~NormalSubtype t1 t2}`;  
 260 that is, is able to decide whether two types are subtypes or not.

261 We will begin by describing and proving correct a version of the algorithm that uses  
 262 explicit type-structural iterators. We will then show the choice stack-based algorithm correct  
 263 by proving equivalence between structural iterators and choice stacks. In doing so, we will  
 264 derive an induction principle for structural iterators (and, as an extension, for choice stacks).

### 3.1 Iterators

The iterator-based implementation is directly equivalent (as will be shown later) to the choice-stack based implementation presented previously in OCaml. However, it retains type structure information inside of the iterator state.

```

269 Inductive TypeIterator: type -> Set :=
270 | TIAtom : forall i, TypeIterator (atom i)
271 | TITuple : forall t1 t2, TypeIterator t1 -> TypeIterator t2 -> TypeIterator (tuple t1 t2)
272 | TIUnionL : forall t1 t2, TypeIterator t1 -> TypeIterator (union t1 t2)
273 | TIUnionR : forall t1 t2, TypeIterator t2 -> TypeIterator (union t1 t2).

```

The `TypeIterator` structure follows the structure of the type being iterated over. Choices at unions are represented as either an instance of `TIUnionR` or `TIUnionL`. This structure then lets us trivially define a function that extracts the current type at the iterator's position:

```

277 Fixpoint current (t:type)(ti:TypeIterator t):type :=
278 match ti with
279 | TIAtom i => atom i
280 | TITuple ti1 ti2 p1 p2 => tuple (current ti1 p1) (current ti2 p2)
281 | TIUnionL ti1 ti2 pl => (current ti1 pl)
282 | TIUnionR ti1 ti2 pr => (current ti2 pr)
283 end.

```

We can then define a function that produces the initial iterator state for a given type:

```

285 Fixpoint start_iterator (t:type):TypeIterator t :=
286 match t with
287 | (atom i) => TIAtom i
288 | (tuple t1 t2) => TITuple t1 t2 (start_iterator t1) (start_iterator t2)
289 | (union t1 t2) => TIUnionL t1 t2 (start_iterator t1)
290 end.

```

Next, we can define a step function that takes one state and either steps it to the next state or indicates that no such next state exists.

```

293 Fixpoint next_state (t:type)(ti:TypeIterator t) : option (TypeIterator t) :=
294 match ti with
295 | TIAtom i => None
296 | TITuple ti1 ti2 p1 p2 =>
297 match (next_state ti2 p2) with
298 | Some np2 => Some(TITuple ti1 ti2 p1 np2)
299 | None =>
300 match (next_state ti1 p1) with
301 | Some np1 => Some(TITuple ti1 ti2 np1 (start_iterator ti2))
302 | None => None
303 end
304 end
305 | TIUnionL ti1 ti2 pl =>
306 match (next_state ti1 pl) with
307 | Some np1 => Some(TIUnionL ti1 ti2 np1)
308 | None => Some(TIUnionR ti1 ti2 (start_iterator ti2))
309 end
310 | TIUnionR ti1 ti2 pr => option_map (TIUnionR ti1 ti2) (next_state ti2 pr)
311 end.

```



312 With these definitions, we can then prove a basic form of correctness with respect to the  
 313 denotational or normalization semantics:

314 ► **Theorem 2** (Correctness of iterators). *Remaining*  $t$  (*start\_iterator*  $t$ ) (*clauses*  $t$ )  
 315 *Every type in  $\llbracket t \rrbracket$  will be explored using *next\_step* from *start\_iterator*  $t$ .*

316 **Proof.** The *Remaining* predicate relates iterators to the list of types that remain to be  
 317 iterated, so the Coq theorem statement indicates that the initial state of the iterator for type  
 318  $t$  has every clause in the normalized version of  $t$  remaining to be iterated.

319 We proceed by induction on  $t$ . The cases for atomic types and unions follow from the IH  
 320 trivially. We prove the theorem for tuples correct by case analyzing on the number of clauses  
 321 induced by the first element in the tuple, then identifying the next element produced by the  
 322 iterator from the tuple.

323 See *iterator\_has\_clauses* in the Coq proof for full details. ◀

324 *next\_state* returns *Some*  $s$  if there is some successor state  $s$  to the current, and *None* if  
 325 the given iterator state is terminal. It will go left-to-right through unions, and will explore  
 326 2-tuples by iterating through the choices on the right for each choice on the left. We can  
 327 then define an induction principle for type iterators based on *next\_state*:

328 ► **Theorem 3.** *Definition iter\_rect*  
 329  $(t : \text{type}) (P : \text{TypeIterator } t \rightarrow \text{Set})$   
 330  $(\text{pi} : \text{forall } it, \text{next\_state } t \text{ it} = \text{None} \rightarrow P \text{ it})$   
 331  $(\text{ps} : \text{forall } it' \text{ it}'', P \text{ it}'' \rightarrow \text{next\_state } t \text{ it}' = \text{Some } it'' \rightarrow P \text{ it}')$   
 332  $(it : \text{TypeIterator } t) : P \text{ it}$   
 333 *For any type  $t$  and proposition  $P$ , and if:*  
 334 ■  *$P$  holds for an iterator that has no next state (e.g. is done)*  
 335 ■ *if  $P$  holds for the following iterator state  $it$ , then  $P$  holds for the preceeding iterator state*  
 336  *$it'$ .*  
 337 *Then  $P$  holds for all iterators for type  $t$*

338 **Proof.** Proving the induction principle for type iterators relies on the *iternum* function,  
 339 which decides the number of steps remaining in the iterator before termination. The proof  
 340 proceeds by simultaneous case analysis on the number of remaining states and whether the  
 341 iterator step function can produce a successor state from the present state.

342 If the iteration number is not yet 0, and if there is a successor state, then we simply  
 343 appeal to the induction hypothesis and continue on. If there is no successor state but the  
 344 iteration number is nonzero or vice versa, then by lemma (*iternum\_monotonic*, taking an  
 345 iterator step decrements the iteration number) contradiction. Finally, if there is no next step  
 346 and the iteration number is 0, then we have reached the base case and terminate.

347 For full details, see the Coq definition of *iter\_rect*. ◀

348 Using *iter\_rect*, we can implement and prove correct equivalent functions to *ex\_subtype*,  
 349 *fa\_ex\_subtype*, and *subtype* as described in the OCaml implementation.

350 **Definition exists\_iter**( $a \ b : \text{type}$ ) :  
 351  $(\{ t \mid \text{InType } t \ b \wedge \text{BaseSubtype } a \ t \} +$   
 352  $\{\text{forall } t, \text{InType } t \ b \rightarrow \sim(\text{BaseSubtype } a \ t)\})).$

353 *exists\_iter* is equivalent to the choice-stack based *ex\_subtype*, and determines if there  
 354 exists some denotationally-contained type in  $b$  that is a supertype of the given  $a$ . Internally, it  
 355 is implemented in the same way as *ex\_subtype*, though using *iter\_rect* to iterate through  
 356 every iterator state.

```

357 Definition forall_iter (a b : type) :
358   { forall t, In t (clauses a) -> exists t', InType t' b /\ (BaseSubtype t t')} +
359   { exists t, In t (clauses a) /\ forall t', InType t' b -> ~ (BaseSubtype t t')}.

```

360 `forall_iter` is to `fa_ex_subtype` what `exists_iter` is to `ex_subtype`. Like `exists_iter` it implements the same decision procedure as `fa_ex_subtype` (and internally relies upon `exists_iter`), though through the abstraction of `iter_rect`.

363 Finally, we can define a decidable function (called `subtype` in the proof) that decides whether two types are subtypes or not. `subtype` simply invokes `forall_iter` to decide subtyping.

```

366 Definition subtype(a b:type) : {NormalSubtype a b} + {~NormalSubtype a b}.
367   destruct (forall_iter a b).
368   - left. [...]
369   - right. [...]
370 Defined.

```

371 Therefore, using iterators, we can decide whether subtyping holds for any two types in our language. We will now show an equivalence between iterators and stacks-of-choices, allowing for more efficient implementation.

## 374 3.2 Stacks

375 To show that the choice-stack based algorithm is correct, we reduce it to the already-shown-correct iterator-based algorithm for deciding subtyping. To do so, we show an equivalence between choice stacks and iterators, then prove correctness of the subtyping algorithms.

378 In the context of the Coq proof, we use the type `st_context` to refer to a choice stack. In Coq, this is represented as a list of boolean values, with false representing a left choice and true representing a right choice at a specific union.

381 To show equivalence between the iterator-based and stack-based algorithm, we need to first prove two properties:

- 383 ■ iterators are convertible to equivalent choice lists;
- 384 ■ stepping an iterator is equivalent to stepping a choice list.

385 We define an iterator and a choice stack to be equivalent if, when applied to the same type, they select the same subset of that type. To describe this, we define `lookup_path` which looks up what type is selected by a given choice stack.

```

388 Fixpoint lookup_path(t:type)(p:st_context) : type * st_context :=
389   match t, p with
390   | atom i, _ => (t, p)
391   | tuple t1 t2, _ =>
392     let (r1,p1) := lookup_path t1 p in
393     let (r2,p2) := lookup_path t2 p1 in
394     (tuple r1 r2, p2)
395   | union l r, false::rs => lookup_path l rs
396   | union l r, true::rs => lookup_path r rs
397   | _, nil => (t, nil)
398   end.

```

399 `lookup_path` is notable in that it both returns the selected type as well as whatever of the choice stack remains once it reaches a leaf. This is needed in order to be able to traverse

types that contain tuples, whose left branches will potentially be given a longer choice stack then necessary.

Next, we can convert iterators to paths in the straightforward manner, as implemented by `iterator_to_path`:

```

405 Fixpoint iterator_to_path(t:type)(it:TypeIterator t):st_context :=
406   match it with
407   | TIAtom _ => nil
408   | TITuple t1 t2 it1 it2 => (iterator_to_path t1 it1) ++ (iterator_to_path t2 it2)
409   | TIUnionL t1 _ it1 => false :: (iterator_to_path t1 it1)
410   | TIUnionR _ t2 it1 => true :: (iterator_to_path t2 it1)
411   end.

```

`iterator_to_path` simply traverses the iterator in order, appending onto the output choice stack whatever choice the iterator makes at that union. This illustrates the equivalence between iterators and choice stacks; choice stacks are simply iterators with the structural information removed.

Using the combination of `lookup_path` and `iterator_to_path`, we can then show the first correctness property that we need to prove that the algorithm using choice stacks is correct:

► **Lemma 4** (Iterator to path is correct). *Lemma `itp_correct` : forall t it, current t it = fst (lookup\_path t (iterator\_to\_path t it)). For every type t and type iterator it, the iterator's current type current t it is equal to the result of looking up the conversion of it to a choice stack.*

**Proof.** See `itp_correct` in the Coq proof. ◀

Stepping in the Coq implementation is implemented identically to the OCaml implementation. It only remains to show that this step operation (called `step_ctx` in Coq) is correct with respect to the iterator `next_state`.

► **Lemma 5** (Correctness of `step_ctx`). *forall t it, step\_ctx t (iterator\_to\_path t it) = (option\_map (iterator\_to\_path t) (next\_state t it)). For every type t and type iterator it, stepping the choice-list equivalent of it will produce the same result as converting the result of stepping it.*

**Proof.** See `list_step_correct` in the Coq proof. ◀

Now, with the relevant properties proven, we can implement and prove correct `ex_subtype` and `fa_ex_subtype` in Coq. The function names are the same, as are the implementations up to the addition of a fuel parameter (which is shown to be unnecessary).

► **Lemma 6** (Correctness of existential subtype checking with choice stacks). *forall a b it, (exists pf, exists\_iter\_inner a b it = inleft pf) <-> exists n, ex\_subtype a b (iterator\_to\_path b it) n = Some true. For every two types a and b, the iterator-based algorithm exists\_iter\_inner will produce a proof that a is a subtype of b if and only if there is an integer n such that ex\_subtype given n fuel runs producing true.*

**Proof.** See `ex_sub_corr_eq` in the Coq proof. ◀

443 ► **Lemma 7** (Correctness of forall-exists subtype checking with choice stacks). forall a b it,  
 444 (exists pf, forall\_iter\_inner a b it = left pf) <->  
 445 exists n, fa\_ex\_subtype a b (iterator\_to\_path a it) n = Some true.

446 *For every two types a and b, the iterator-based algorithm forall\_iter\_inner will produce*  
 447 *a proof that a is a subtype of b if and only if there is an integer n such that fa\_ex\_subtype*  
 448 *given n fuel runs producing true.*

449 **Proof.** See fa\_sub\_corr\_eq in the Coq proof. ◀

450 The choice stack-based algorithm therefore is provably equivalent to the iterator-based  
 451 algorithm, and is thus correct.

## 452 **4 Performance Analysis**

453 The algorithm improves upon normalization in two key ways:

- 454 ■ lazily exploring possible clauses, obviating the need to store a fully normalized type;
- 455 ■ enabling fast paths that avoid the exploration of the full choice space.

456 In the worst case, our algorithm has the same big-O time complexity. However, lazily  
 457 exploring the choice space allows us to require worst-case polynomial space, in comparison to  
 458 normalization's exponential space complexity. Similarly, the algorithm enables optimizations  
 459 that offer best (and typical) case time complexity improvements from exponential to linear  
 460 time.

461 Worst case time complexity of both subtyping algorithms is determined by the number of  
 462 clauses that would exist in the normalized type. In the worst case, (a tuple of unions), each  
 463 union begets a different clause in the normalized type. Consider `Tuple{Union{A, B}, Union{C, D}}`,  
 464 which will normalize to `Union{Tuple{A, C}, Tuple{A, D}, Tuple{B, C}, Tuple{B, D}}` gen-  
 465 erating a new tuple for each choice for every contained union. As a result, there are worst-case  
 466  $2^n$  tuples in the fully normalized version of a type that has  $n$  unions.

467 In order to ensure correctness, each of these tuples (or choices at unions) must always  
 468 be explored. As a result, both the algorithm we present here and normalization will have  
 469 worst-case  $O(2^n)$  time complexity. The approaches differ, however, in space complexity. The  
 470 normalization approach computes and stores each of the exponentially many alternatives,  
 471 so also has  $O(2^n)$  space complexity. However, the algorithm we discuss need only store the  
 472 choice made at each union, thereby offering  $O(n)$  space complexity.

473 The algorithm we discuss also can improve best-case time performance. Normalization  
 474 will necessarily be  $o(2^n)$  due to computation of the entire normalized type. However, the lazy  
 475 subtyping algorithm need only make one choice before discovering that a subtype relation  
 476 exists in the best-case, giving  $o(n)$  performance. Moreover, computing type choices lazily  
 477 enables fast-paths to short circuit full exploration of choice alternatives.

478 This is important for Julia due to a common programming idiom. Many Julia library  
 479 developers write signatures of the form `Tuple{Union{A, B}, Union{C, D}}` to indicate that  
 480 their method can take any of the named types. When deciding dispatch against these  
 481 methods, Julia will frequently check if a tuple containing solely concrete (instantiable) types  
 482 is a subtype of the tuple of unions. If Julia used normalization, this would always be  
 483 exponential on the number of unions that appeared in the argument list as this is the above  
 484 mentioned worst-case exponential complexity case. However, its use of the lazy algorithm  
 485 enables it frequently identify the best alternative and short circuit before having to explore  
 486 much of the choice possibility space.

## 5 Conclusion

We have presented an algorithm for deciding subtyping relationships between types that consist of atomic types, tuples, and unions. This algorithm is able to decide subtyping relationships in the presence of distributive semantics for union types without needing normalization (and therefore using linear space) and without additionally constraining type system features.

---

## References

- 1 Hack. <https://hacklang.org/>. Accessed: 2019-01-11.
- 2 Jeff Bezanson, Alan Edelman, Stefan Karpinski, and Viral B. Shah. Julia: A fresh approach to numerical computing. *SIAM Review*, 59(1), 2017. doi:10.1137/141000671.
- 3 Alain Frisch, Giuseppe Castagna, and Véronique Benzaken. Semantic subtyping: Dealing set-theoretically with function, union, intersection, and negation types. *J. ACM*, 55(4), 2008. doi:10.1145/1391289.1391293.
- 4 Anders Hejlsberg. Introducing typescript. URL: <https://channel9.msdn.com/posts/Anders-Hejlsberg-Introducing-TypeScript>.
- 5 Benjamin Pierce. Programming with intersection types, union types, and polymorphism. Technical Report CMU-CS-91-106, Carnegie Mellon University, 1991.
- 6 Jerome Vouillon. Subtyping union types. In *Computer Science Logic (CSL)*, 2004. doi:10.1007/978-3-540-30124-0\_32.
- 7 Francesco Zappa Nardelli, Julia Belyakova, Artem Pelenitsyn, Benjamin Chung, Jeff Bezanson, and Jan Vitek. Julia subtyping: a rational reconstruction. *PACMPL*, 2(OOPSLA):113:1–113:27, 2018. URL: <https://doi.org/10.1145/3276483>, doi:10.1145/3276483.